

```

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
# plt.use('Agg')
%matplotlib inline
z = np.random.rand(21315)

data = pd.read_csv(r"C:\Users\kadam\OneDrive\Desktop\jyp_python\miniproject\archive (26)\city_day.csv")
data.to_pickle("data.pkl")

```

Data Processing

```

data.shape
(29531, 16)

data.head()

      City        Date   PM2.5   PM10     NO    NO2    NOx   NH3     CO
S02 \
0 Ahmedabad  2015-01-01      NaN      NaN  0.92  18.22  17.15  NaN  0.92
27.64
1 Ahmedabad  2015-01-02      NaN      NaN  0.97  15.69  16.46  NaN  0.97
24.55
2 Ahmedabad  2015-01-03      NaN      NaN  17.40  19.30  29.70  NaN  17.40
29.07
3 Ahmedabad  2015-01-04      NaN      NaN  1.70  18.48  17.97  NaN  1.70
18.59
4 Ahmedabad  2015-01-05      NaN      NaN  22.10  21.42  37.76  NaN  22.10
39.33

      O3  Benzene  Toluene  Xylene  AQI  AQI_Bucket
0  133.36     0.00     0.02     0.00  NaN      NaN
1   34.06     3.68     5.50     3.77  NaN      NaN
2   30.70     6.80    16.40     2.25  NaN      NaN
3   36.08     4.43    10.14     1.00  NaN      NaN
4   39.31     7.01    18.89     2.78  NaN      NaN

data.tail()

      City        Date   PM2.5   PM10     NO    NO2    NOx
NH3 \
29526 Visakhapatnam  2020-06-27  15.02  50.94  7.68  25.06  19.54
12.47
29527 Visakhapatnam  2020-06-28  24.38  74.09  3.42  26.06  16.53
11.99
29528 Visakhapatnam  2020-06-29  22.91  65.73  3.45  29.53  18.33

```

```
10.71  
29529 Visakhapatnam 2020-06-30 16.64 49.97 4.05 29.26 18.80  
10.03  
29530 Visakhapatnam 2020-07-01 15.00 66.00 0.40 26.85 14.05  
5.20
```

	C0	S02	03	Benzene	Toluene	Xylene	AQI
AQI_Bucket							
29526 Good	0.47	8.55	23.30	2.24	12.07	0.73	41.0
29527 Satisfactory	0.52	12.72	30.14	0.74	2.21	0.38	70.0
29528 Satisfactory	0.48	8.42	30.96	0.01	0.01	0.00	68.0
29529 Satisfactory	0.52	9.84	28.30	0.00	0.00	0.00	54.0
29530 Good	0.59	2.10	17.05	NaN	NaN	NaN	50.0

```
data.dtypes
```

```
City          object  
Date         object  
PM2.5        float64  
PM10         float64  
NO           float64  
NO2          float64  
NOx          float64  
NH3          float64  
C0           float64  
S02          float64  
03           float64  
Benzene      float64  
Toluene      float64  
Xylene       float64  
AQI          float64  
AQI_Bucket   object  
dtype: object
```

```
data.columns
```

```
Index(['City', 'Date', 'PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3',  
'C0', 'S02',  
       '03', 'Benzene', 'Toluene', 'Xylene', 'AQI', 'AQI_Bucket'],  
      dtype='object')
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 29531 entries, 0 to 29530  
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	City	29531	non-null object
1	Date	29531	non-null object
2	PM2.5	24933	non-null float64
3	PM10	18391	non-null float64
4	NO	25949	non-null float64
5	NO2	25946	non-null float64
6	NOx	25346	non-null float64
7	NH3	19203	non-null float64
8	CO	27472	non-null float64
9	S02	25677	non-null float64
10	O3	25509	non-null float64
11	Benzene	23908	non-null float64
12	Toluene	21490	non-null float64
13	Xylene	11422	non-null float64
14	AQI	24850	non-null float64
15	AQI_Bucket	24850	non-null object

dtypes: float64(13), object(3)
memory usage: 3.6+ MB

dataset have null values.

It doesn't have invalid datatypes.

```
data_null = np.where(data.isnull() == True)
data_null

(array([ 0, 0, 0, ..., 29530, 29530, 29530]),
 array([ 2, 3, 7, ..., 11, 12, 13]))
```


data.describe()

	PM2.5	PM10	NO	NO2
NOx \ count	24933.000000	18391.000000	25949.000000	25946.000000
25346.000000				
mean	67.450578	118.127103	17.574730	28.560659
32.309123				
std	64.661449	90.605110	22.785846	24.474746
31.646011				
min	0.040000	0.010000	0.020000	0.010000
0.000000				
25%	28.820000	56.255000	5.630000	11.750000
12.820000				
50%	48.570000	95.680000	9.890000	21.690000
23.520000				
75%	80.590000	149.745000	19.950000	37.620000
40.127500				
max	949.990000	1000.000000	390.680000	362.210000
467.630000				


	NH3	CO	SO2	03
Benzene \				
count	19203.000000	27472.000000	25677.000000	25509.000000
	23908.000000			
mean	23.483476	2.248598	14.531977	34.491430
	3.280840			
std	25.684275	6.962884	18.133775	21.694928
	15.811136			
min	0.010000	0.000000	0.010000	0.010000
	0.000000			
25%	8.580000	0.510000	5.670000	18.860000
	0.120000			
50%	15.850000	0.890000	9.160000	30.840000
	1.070000			
75%	30.020000	1.450000	15.220000	45.570000
	3.080000			
max	352.890000	175.810000	193.860000	257.730000
	455.030000			
	Toluene	Xylene	AQI	
count	21490.000000	11422.000000	24850.000000	
mean	8.700972	3.070128	166.463581	
std	19.969164	6.323247	140.696585	
min	0.000000	0.000000	13.000000	
25%	0.600000	0.140000	81.000000	
50%	2.970000	0.980000	118.000000	
75%	9.150000	3.350000	208.000000	
max	454.850000	170.370000	2049.000000	

Data Cleaning

```
data['AQI'].plot(kind='hist')
plt.show()
```




```
data['AQI'].plot(kind='bar')
plt.show()
```



```
data.boxplot()
```

```
<Axes: >
```



```
data['PM2.5'].fillna(data['PM2.5'].mean(), inplace=True)
#df['column_name'].fillna(df['column_name'].mean(), inplace=True)
```

C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\351580854.py:1:
 FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
 The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
data['PM2.5'].fillna(data['PM2.5'].mean(), inplace=True)
```

data

	City	Date	PM2.5	PM10	NO	NO2
NOx \ 0	Ahmedabad	2015-01-01	67.450578	NaN	0.92	18.22
17.15	Ahmedabad	2015-01-02	67.450578	NaN	0.97	15.69
16.46						


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29531 entries, 0 to 29530
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   City         29531 non-null   object  
 1   Date          29531 non-null   object  
 2   PM2.5        29531 non-null   float64 
 3   PM10         18391 non-null   float64 
 4   NO            25949 non-null   float64 
 5   NO2           25946 non-null   float64 
 6   NOx          25346 non-null   float64 
 7   NH3           19203 non-null   float64 
 8   CO            27472 non-null   float64 
 9   SO2           25677 non-null   float64 
 10  O3            25509 non-null   float64 
 11  Benzene       23908 non-null   float64 
 12  Toluene       21490 non-null   float64 
 13  Xylene        11422 non-null   float64 
 14  AQI           24850 non-null   float64 
 15  AQI_Bucket    24850 non-null   object  
dtypes: float64(13), object(3)
memory usage: 3.6+ MB
```

```
data['PM2.5'].fillna(data['PM2.5'].mean(), inplace=True)
```

```
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\935467664.py:1:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
data['PM2.5'].fillna(data['PM2.5'].mean(), inplace=True)
```

```
data['PM10'].fillna(data['PM10'].mean(), inplace=True)
```

```
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\715378189.py:1:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try

```
using 'df.method({col: value}, inplace=True)' or df[col] =  
df[col].method(value) instead, to perform the operation inplace on the  
original object.
```

```
data['PM10'].fillna(data['PM10'].mean(), inplace=True)  
data['NO'].fillna(data['NO'].mean(), inplace=True)
```

```
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\3552883858.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or  
Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never  
work because the intermediate object on which we are setting values  
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

```
data['NO'].fillna(data['NO'].mean(), inplace=True)  
data['NO2'].fillna(data['NO2'].mean(), inplace=True)
```

```
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\2390402205.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or  
Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never  
work because the intermediate object on which we are setting values  
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

```
data['NO2'].fillna(data['NO2'].mean(), inplace=True)  
data['NOx'].fillna(data['NOx'].mean(), inplace=True)
```

```
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\2133220639.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or  
Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never  
work because the intermediate object on which we are setting values  
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try

```
using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.
```

```
data['NOx'].fillna(data['NOx'].mean(), inplace=True)  
data['NH3'].fillna(data['NH3'].mean(), inplace=True)  
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\2854111976.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
data['NH3'].fillna(data['NH3'].mean(), inplace=True)  
data['CO'].fillna(data['CO'].mean(), inplace=True)  
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\1161907318.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
data['CO'].fillna(data['CO'].mean(), inplace=True)  
data['SO2'].fillna(data['SO2'].mean(), inplace=True)  
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\3344645610.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try

```
using 'df.method({col: value}, inplace=True)' or df[col] =  
df[col].method(value) instead, to perform the operation inplace on the  
original object.
```

```
data['S02'].fillna(data['S02'].mean(), inplace=True)  
data['03'].fillna(data['03'].mean(), inplace=True)  
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\2730838331.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or  
Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never  
work because the intermediate object on which we are setting values  
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

```
data['03'].fillna(data['03'].mean(), inplace=True)  
data['Benzene'].fillna(data['Benzene'].mean(), inplace=True)  
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\4277266277.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or  
Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never  
work because the intermediate object on which we are setting values  
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

```
data['Benzene'].fillna(data['Benzene'].mean(), inplace=True)  
data['Toluene'].fillna(data['Toluene'].mean(), inplace=True)  
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\4248952095.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or  
Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never  
work because the intermediate object on which we are setting values  
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try

```
using 'df.method({col: value}, inplace=True)' or df[col] =  
df[col].method(value) instead, to perform the operation inplace on the  
original object.
```

```
data['Toluene'].fillna(data['Toluene'].mean(), inplace=True)  
data['Xylene'].fillna(data['Xylene'].mean(), inplace=True)  
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\2593633129.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or  
Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never  
work because the intermediate object on which we are setting values  
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

```
data['Xylene'].fillna(data['Xylene'].mean(), inplace=True)  
data['AQI'].fillna(data['AQI'].mean(), inplace=True)  
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\2966707500.py:1:  
FutureWarning: A value is trying to be set on a copy of a DataFrame or  
Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never  
work because the intermediate object on which we are setting values  
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

```
data['AQI'].fillna(data['AQI'].mean(), inplace=True)  
data.describe()  


|              | PM2.5        | PM10         | NO           | NO2          |
|--------------|--------------|--------------|--------------|--------------|
| NOx \        |              |              |              |              |
| count        | 29531.000000 | 29531.000000 | 29531.000000 | 29531.000000 |
| 29531.000000 |              |              |              |              |
| mean         | 67.450578    | 118.127103   | 17.57473     | 28.560659    |
| 32.309123    |              |              |              |              |
| std          | 59.414476    | 71.500953    | 21.35922     | 22.941051    |
| 29.317936    |              |              |              |              |


```

```
min      0.040000      0.010000      0.02000      0.010000  
0.000000  
25%    32.150000    79.315000     6.21000    12.980000  
14.670000  
50%    58.030000   118.127103    11.53000    25.240000  
27.550000  
75%    72.450000   118.127103    17.57473    34.665000  
36.015000  
max    949.990000  1000.000000    390.68000   362.210000  
467.630000
```

	NH3	CO	SO2	O3
Benzene \				
count	29531.000000	29531.000000	29531.000000	29531.000000
29531.000000				
mean	23.483476	2.248598	14.531977	34.491430
3.280840				
std	20.711370	6.715753	16.909088	20.163443
14.226364				
min	0.010000	0.000000	0.010000	0.010000
0.000000				
25%	12.040000	0.540000	6.090000	20.740000
0.240000				
50%	23.483476	0.950000	10.480000	34.491430
1.840000				
75%	23.483476	1.710000	14.531977	42.730000
3.280840				
max	352.890000	175.810000	193.860000	257.730000
455.030000				

	Toluene	Xylene	AQI
count	29531.000000	29531.000000	29531.000000
mean	8.700972	3.070128	166.463581
std	17.034769	3.932426	129.064348
min	0.000000	0.000000	13.000000
25%	1.280000	2.000000	88.000000
50%	6.930000	3.070128	138.000000
75%	8.700972	3.070128	179.000000
max	454.850000	170.370000	2049.000000


```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 29531 entries, 0 to 29530  
Data columns (total 16 columns):  
 #   Column      Non-Null Count  Dtype     
---    
 0   City        29531 non-null   object    
 1   Date        29531 non-null   object    
 2   PM2.5       29531 non-null   float64
```

```
3    PM10        29531 non-null  float64
4    NO          29531 non-null  float64
5    NO2         29531 non-null  float64
6    NOx         29531 non-null  float64
7    NH3          29531 non-null  float64
8    CO          29531 non-null  float64
9    SO2          29531 non-null  float64
10   O3          29531 non-null  float64
11   Benzene      29531 non-null  float64
12   Toluene      29531 non-null  float64
13   Xylene        29531 non-null  float64
14   AQI          29531 non-null  float64
15   AQI_Bucket    24850 non-null  object
dtypes: float64(13), object(3)
memory usage: 3.6+ MB
```

```
plt.figure(figsize=(12,6))
data.boxplot()
```

<Axes: >




```
city = data.City.value_counts()
city

City
Ahmedabad      2009
Bengaluru       2009
Chennai         2009
Mumbai          2009
Lucknow         2009
```

```
Delhi          2009
Hyderabad      2006
Patna          1858
Gurugram       1679
Visakhapatnam 1462
Amritsar        1221
Jorapokhar     1169
Jaipur          1114
Thiruvananthapuram 1112
Amaravati      951
Brajrajnagar    938
Talcher         925
Kolkata         814
Guwahati        502
Coimbatore      386
Shillong         310
Chandigarh      304
Bhopal           289
Kochi            162
Ernakulam        162
Aizawl           113
Name: count, dtype: int64
```


```
plt.figure(figsize=(12,6))
datahm = data[(data['City']=='Ahmedabad')]
datahm.boxplot()
```

```
<Axes: >
```



```
datahm.describe()
```

	PM2.5	PM10	NO	NO2	NOx
\count	2009.000000	2009.000000	2009.000000	2009.000000	2009.000000
mean	67.728234	117.409318	20.956815	49.805675	42.914773
std	32.739185	20.608768	18.030998	41.889674	29.023976
min	3.040000	11.500000	0.060000	0.080000	0.000000
25%	46.910000	118.127103	10.380000	28.560659	27.840000
50%	67.450578	118.127103	17.574730	28.560659	32.309123
75%	73.070000	118.127103	23.750000	66.430000	51.030000
max	381.690000	586.270000	175.810000	292.020000	246.030000
NH3 CO SO2 O3					
Benzene \count	2.009000e+03	2009.000000	2009.000000	2009.000000	2009.000000
mean	2.348348e+01	16.147420	42.281148	37.565152	4.901003
std	7.107196e-15	20.258113	37.926831	18.464239	6.953368
min	2.348348e+01	0.060000	0.520000	0.380000	0.000000
25%	2.348348e+01	2.248598	14.531977	32.100000	1.820000
50%	2.348348e+01	8.510000	23.810000	34.491430	3.280840
75%	2.348348e+01	23.750000	60.680000	45.650000	4.720000
max	2.348348e+01	175.810000	186.080000	162.430000	115.140000
Toluene Xylene AQI					
count	2009.000000	2009.000000	2009.000000		
mean	23.163071	3.964491	356.144807		
std	26.787328	6.547374	287.617151		
min	0.000000	0.000000	48.000000		
25%	8.700972	0.660000	166.463581		
50%	11.320000	3.070128	229.000000		
75%	32.330000	3.850000	465.000000		
max	371.650000	109.230000	2049.000000		
datahm.boxplot(column='PM10')					
<Axes: >					



```

col =
['PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene', 'Xylene', 'AQI']

Q3 = datahm[col].quantile(0.75)
Q1 = datahm[col].quantile(0.25)

Q1,Q3

(PM2.5      46.910000
PM10       118.127103
NO         10.380000
NO2        28.560659
NOx        27.840000
NH3        23.483476
CO          2.248598
SO2        14.531977
O3         32.100000
Benzene    1.820000
Toluene    8.700972
Xylene     0.660000
AQI        166.463581
Name: 0.25, dtype: float64,
PM2.5      73.070000
PM10       118.127103
NO         23.750000

```

```
N02      66.430000  
N0x      51.030000  
NH3      23.483476  
C0       23.750000  
S02      60.680000  
O3       45.650000  
Benzene   4.720000  
Toluene   32.330000  
Xylene    3.850000  
AQI      465.000000  
Name: 0.75, dtype: float64)
```

```
IQR = Q3 - Q1
```

```
IQR
```

```
PM2.5     26.160000  
PM10      0.000000  
NO        13.370000  
N02      37.869341  
N0x      23.190000  
NH3      0.000000  
C0       21.501402  
S02      46.148023  
O3       13.550000  
Benzene   2.900000  
Toluene   23.629028  
Xylene    3.190000  
AQI      298.536419  
dtype: float64
```

```
lower_limit = Q1 - 1.5*IQR  
upper_limit = Q3 + 1.5*IQR  
lower_limit, upper_limit
```

```
(PM2.5      7.670000  
PM10      118.127103  
NO        -9.675000  
N02      -28.243352  
N0x      -6.945000  
NH3      23.483476  
C0       -30.003504  
S02      -54.690057  
O3       11.775000  
Benzene   -2.530000  
Toluene   -26.742570  
Xylene    -4.125000  
AQI      -281.341046  
dtype: float64,  
PM2.5      112.310000  
PM10      118.127103
```

```

NO          43.805000
N02         123.234011
N0x         85.815000
NH3          23.483476
CO           56.002103
S02         129.902034
O3           65.975000
Benzene      9.070000
Toluene       67.773542
Xylene        8.635000
AQI          912.804628
dtype: float64)

```

```
datahm[(datahm[col]<lower_limit)|(datahm[col]>upper_limit)]
```

	City	Date	PM2.5	PM10	NO	N02	N0x	NH3	CO	S02	O3
Benzene	\										
0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	133.36
NaN											
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN											
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN											
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN											
4	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN											
...
...											
2004	NaN	NaN	NaN	118.67	NaN	NaN	NaN	NaN	NaN	NaN	68.05
NaN											
2005	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN											
2006	NaN	NaN	NaN	127.98	NaN						
NaN											
2007	NaN	NaN	NaN	121.10	NaN						
NaN											
2008	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	9.69
NaN											

	Toluene	Xylene	AQI	AQI_Bucket
0	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN
...
2004	NaN	NaN	NaN	NaN
2005	NaN	NaN	NaN	NaN
2006	NaN	NaN	NaN	NaN

2007	NaN	NaN	NaN	NaN
2008	NaN	NaN	NaN	NaN

[2009 rows x 16 columns]

```
datahmo = datahm[(datahm[col]>lower_limit) &
(datahm[col]<upper_limit)]
datahmo
```


	City	Date	PM2.5	PM10	NO	N02	N0x	NH3	CO
S02	\								
0	NaN	NaN	67.450578	NaN	0.92	18.22	17.15	NaN	0.92
27.64									
1	NaN	NaN	67.450578	NaN	0.97	15.69	16.46	NaN	0.97
24.55									
2	NaN	NaN	67.450578	NaN	17.40	19.30	29.70	NaN	17.40
29.07									
3	NaN	NaN	67.450578	NaN	1.70	18.48	17.97	NaN	1.70
18.59									
4	NaN	NaN	67.450578	NaN	22.10	21.42	37.76	NaN	22.10
39.33									
...
.									
2004	NaN	NaN	62.120000	NaN	9.18	56.35	19.86	NaN	0.49
12.44									
2005	NaN	NaN	31.570000	NaN	6.37	23.99	16.40	NaN	0.52
11.01									
2006	NaN	NaN	29.750000	NaN	9.06	25.15	18.92	NaN	0.67
12.10									
2007	NaN	NaN	40.020000	NaN	7.09	58.92	33.41	NaN	0.73
16.39									
2008	NaN	NaN	37.630000	NaN	4.42	35.04	20.17	NaN	0.28
14.40									

	03	Benzene	Toluene	Xylene	AQI	AQI_Bucket
0	NaN	0.00	0.02	0.00	166.463581	NaN
1	34.06	3.68	5.50	3.77	166.463581	NaN
2	30.70	6.80	16.40	2.25	166.463581	NaN
3	36.08	4.43	10.14	1.00	166.463581	NaN
4	39.31	7.01	18.89	2.78	166.463581	NaN
...
2004	NaN	1.32	37.76	1.62	92.000000	NaN
2005	26.34	1.37	49.58	1.34	82.000000	NaN
2006	34.99	1.39	60.21	0.79	74.000000	NaN
2007	41.64	1.21	44.10	1.35	98.000000	NaN
2008	NaN	1.73	47.05	1.87	119.000000	NaN

[2009 rows x 16 columns]

```
plt.figure(figsize=(12,6))
datahmo.boxplot()
```

```
<Axes: >
```



```
data.isnull().sum()
```

```
City          0
Date          0
PM2.5         0
PM10          0
NO            0
NO2           0
NOx           0
NH3           0
CO            0
SO2           0
O3            0
Benzene       0
Toluene       0
Xylene        0
AQI          0
AQI_Bucket    4681
dtype: int64
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29531 entries, 0 to 29530
Data columns (total 16 columns):
```

```
#   Column      Non-Null Count   Dtype  
--- 
0   City          29531 non-null    object  
1   Date          29531 non-null    object  
2   PM2.5         29531 non-null    float64 
3   PM10          29531 non-null    float64 
4   NO            29531 non-null    float64 
5   NO2           29531 non-null    float64 
6   NOx           29531 non-null    float64 
7   NH3           29531 non-null    float64 
8   CO            29531 non-null    float64 
9   SO2           29531 non-null    float64 
10  O3            29531 non-null    float64 
11  Benzene       29531 non-null    float64 
12  Toluene       29531 non-null    float64 
13  Xylene        29531 non-null    float64 
14  AQI           29531 non-null    float64 
15  AQI_Bucket    24850 non-null    object  
dtypes: float64(13), object(3)
memory usage: 3.6+ MB
```

```
dataac = data.dropna()

dataac.to_csv('clean_city.csv', index=False)

cleandt = pd.read_csv(r'clean_city.csv')
cleandt.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24850 entries, 0 to 24849
Data columns (total 16 columns):
 #   Column      Non-Null Count   Dtype  
--- 
0   City          24850 non-null    object  
1   Date          24850 non-null    object  
2   PM2.5         24850 non-null    float64 
3   PM10          24850 non-null    float64 
4   NO            24850 non-null    float64 
5   NO2           24850 non-null    float64 
6   NOx           24850 non-null    float64 
7   NH3           24850 non-null    float64 
8   CO            24850 non-null    float64 
9   SO2           24850 non-null    float64 
10  O3            24850 non-null    float64 
11  Benzene       24850 non-null    float64 
12  Toluene       24850 non-null    float64 
13  Xylene        24850 non-null    float64 
14  AQI           24850 non-null    float64 
15  AQI_Bucket    24850 non-null    object
```

```
dtypes: float64(13), object(3)
memory usage: 3.0+ MB

print(data["AQI_Bucket"].value_counts().sum())
print(cleandt["AQI_Bucket"].value_counts().sum())

24850
24850

data['AQI_Bucket'].unique()

array([nan, 'Poor', 'Very Poor', 'Severe', 'Moderate', 'Satisfactory',
       'Good'], dtype=object)

data['AQI_Bucket'].nunique()

6

data['AQI_Bucket'].value_counts()

AQI_Bucket
Moderate      8829
Satisfactory  8224
Poor          2781
Very Poor     2337
Good          1341
Severe         1338
Name: count, dtype: int64

data = data.drop_duplicates()


data.shape

(29531, 16)


cleandt = cleandt.drop_duplicates()
cleandt.shape

(24850, 16)

sns.countplot(cleandt['AQI_Bucket'], color='magenta')
plt.show()
```




```
cleandt['AQI_Bucket'].value_counts().plot(kind='bar', color='green')
plt.show()
```




```
#how PM10 is affecting overall aqi
sns.scatterplot(x=cleandt['AQI'],
y=cleandt['PM10'],hue=cleandt['AQI_Bucket'])

<Axes: xlabel='AQI', ylabel='PM10'>
```



```
sns.scatterplot(x=cleandt['NH3'],
y=cleandt['NOx'],hue=cleandt['AQI_Bucket'])

<Axes: xlabel='NH3', ylabel='NOx'>
```




Combined effect of NOx and NH3 :

NOx(250) and NH3(100) this combination is really harmful


```
sns.scatterplot(x=cleandt['AQI'],y=cleandt['Benzene'],hue=cleandt['AQI_Bucket'])
```

```
<Axes: xlabel='AQI', ylabel='Benzene'>
```




benzene doesn't impact thoroughly on air quality


```
#categorical and numerical
sns.barplot(x=cleandt["AQI_Bucket"],y=cleandt["AQI"],color='slategrey')
<Axes: xlabel='AQI_Bucket', ylabel='AQI'>
```




```
sns.barplot(x=cleandt["AQI_Bucket"],y=cleandt["NH3"],color='indigo')  
<Axes: xlabel='AQI_Bucket', ylabel='NH3'>
```



```
sns.histplot(data=cleandt,kde=True)
<Axes: ylabel='Count'>
C:\Users\kadam\AppData\Local\Programs\Python\Python312\Lib\site-
packages\IPython\core\events.py:82: UserWarning: Creating legend with
loc="best" can be slow with large amounts of data.
    func(*args, **kwargs)
C:\Users\kadam\AppData\Local\Programs\Python\Python312\Lib\site-
packages\IPython\core\pylabtools.py:152: UserWarning: Creating legend
with loc="best" can be slow with large amounts of data.
    fig.canvas.print_figure(bytes_io, **kw)
```




```
sns.histplot(data=cleandt,x=cleandt[ 'AQI' ],kde=True)  
<Axes: xlabel='AQI', ylabel='Count'>
```




maximum no. of cities have AQI in between 50 to 200

```
relu=pd.crosstab(cleandt['AQI'] , cleandt['AQI_Bucket'])  
sns.heatmap(relu)
```

```
<Axes: xlabel='AQI_Bucket', ylabel='AQI'>
```




```
sns.pairplot(cleandt,hue='AQI_Bucket')  
<seaborn.axisgrid.PairGrid at 0x215239d90d0>
```



```
# Grouping the data by city and calculating the mean AQI for each city
city_aqi = cleandt.groupby('City')
['AQI'].mean().sort_values(ascending=False)

# Matplotlib bar plot
plt.figure(figsize=(10, 6))
plt.bar(city_aqi.index, city_aqi.values, color='skyblue')
plt.xticks(rotation=90)
plt.title('City-wise Mean AQI (Matplotlib)')
plt.ylabel('Mean AQI')
plt.xlabel('City')
plt.show()
```


```
# Seaborn bar plot
plt.figure(figsize=(10, 6))
sns.barplot(x=city_aqi.index, y=city_aqi.values, palette='viridis')
plt.xticks(rotation=90)
plt.title('City-wise Mean AQI (Seaborn)')
plt.ylabel('Mean AQI')
plt.xlabel('City')
plt.show()
```



C:\Users\kadam\AppData\Local\Temp\ipykernel_11664\2716694310.py:15:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=city_aqi.index, y=city_aqi.values, palette='viridis')
```



Ahmedabad has Highest Aqi.

```
# Group data by city to analyze each region's trends
city_grouped_data = cleandt.groupby('City')

# Function to generate region-specific conclusions
def region_specific_conclusions(city_group):
    # Get the average AQI for the city
    avg_aqi = city_group['AQI'].astype(float).mean()

    # Find the most frequent AQI bucket (category)
    most_common_aqi_bucket = city_group['AQI_Bucket'].mode()[0]

    # Get the average pollutant levels (PM2.5, PM10, NO2, CO)
    avg_pm25 = city_group['PM2.5'].mean()
    avg_pm10 = city_group['PM10'].mean()
    avg_no2 = city_group['NO2'].mean()
    avg_co = city_group['CO'].mean()

    # Form conclusions
    conclusion = f"For the city of {city_group['City'].iloc[0]}:\n"
    conclusion += f"Mean AQI: {avg_aqi}\n"
    conclusion += f"Most Common AQI Bucket: {most_common_aqi_bucket}\n"
    conclusion += f"Average Pollutant Levels: PM2.5={avg_pm25}, PM10={avg_pm10}, NO2={avg_no2}, CO={avg_co}\n"
    return conclusion
```

```

        conclusion += f"- The average AQI is {avg_aqi:.2f}, falling into
the '{most_common_aqi_bucket}' category.\n"
        conclusion += f"- Key pollutant levels on average are: PM2.5:
{avg_pm25:.2f}, PM10: {avg_pm10:.2f}, NO2: {avg_no2:.2f}, and CO:
{avg_co:.2f}.\n"

    # Identify a potential concern based on high average pollutant
levels
    if avg_pm25 > 60:
        conclusion += f"- PM2.5 levels are relatively high, which may
pose a health risk.\n"
    if avg_pm10 > 100:
        conclusion += f"- PM10 levels are also elevated, contributing
to reduced air quality.\n"
    if avg_no2 > 40:
        conclusion += f"- NO2 levels exceed safe limits, indicating a
concern for air pollution from combustion sources.\n"

return conclusion

# Generate conclusions for each city
city_conclusions = {}
for city, group in city_grouped_data:
    city_conclusions[city] = region_specific_conclusions(group)

# Display the conclusions for each city
for city, conclusion in city_conclusions.items():
    print(conclusion)

# Visualization of AQI and Pollutants per city

# Convert AQI to numeric for analysis
cleandt['AQI'] = pd.to_numeric(cleandt['AQI'], errors='coerce')

# Plotting average AQI for each city
plt.figure(figsize=(10,6))
sns.barplot(x='City', y='AQI', data=cleandt, estimator=lambda x:
x.mean())
plt.title('Average AQI per City')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Plotting average PM2.5 and PM10 levels per city
plt.figure(figsize=(12,6))
sns.barplot(x='City', y='PM2.5', data=cleandt, estimator=lambda x:
x.mean(), color='blue', label='PM2.5')
sns.barplot(x='City', y='PM10', data=cleandt, estimator=lambda x:
x.mean(), color='orange', label='PM10')
plt.title('Average PM2.5 and PM10 Levels per City')

```

```
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()
```

For the city of Ahmedabad:

- The average AQI is 452.12, falling into the 'Severe' category.
- Key pollutant levels on average are: PM2.5: 67.87, PM10: 116.54, NO₂: 60.18, and CO: 22.09.
- PM2.5 levels are relatively high, which may pose a health risk.
- PM10 levels are also elevated, contributing to reduced air quality.
- NO₂ levels exceed safe limits, indicating a concern for air pollution from combustion sources.

For the city of Aizawl:

- The average AQI is 34.77, falling into the 'Good' category.
- Key pollutant levels on average are: PM2.5: 17.44, PM10: 24.14, NO₂: 0.37, and CO: 0.28.

For the city of Amaravati:

- The average AQI is 95.30, falling into the 'Satisfactory' category.
- Key pollutant levels on average are: PM2.5: 38.25, PM10: 76.21, NO₂: 21.73, and CO: 0.69.

For the city of Amritsar:

- The average AQI is 119.92, falling into the 'Satisfactory' category.
- Key pollutant levels on average are: PM2.5: 56.01, PM10: 115.10, NO₂: 18.54, and CO: 0.62.
- PM10 levels are also elevated, contributing to reduced air quality.

For the city of Bengaluru:

- The average AQI is 94.32, falling into the 'Satisfactory' category.
- Key pollutant levels on average are: PM2.5: 36.78, PM10: 88.39, NO₂: 28.31, and CO: 1.67.

For the city of Bhopal:

- The average AQI is 132.83, falling into the 'Moderate' category.
- Key pollutant levels on average are: PM2.5: 50.21, PM10: 119.72, NO₂: 31.37, and CO: 0.88.
- PM10 levels are also elevated, contributing to reduced air quality.

For the city of Brajrajnagar:

- The average AQI is 150.28, falling into the 'Moderate' category.
- Key pollutant levels on average are: PM2.5: 63.99, PM10: 123.44, NO₂: 18.23, and CO: 1.84.
- PM2.5 levels are relatively high, which may pose a health risk.
- PM10 levels are also elevated, contributing to reduced air quality.

For the city of Chandigarh:

- The average AQI is 96.50, falling into the 'Satisfactory' category.

- Key pollutant levels on average are: PM2.5: 42.36, PM10: 85.80, NO₂: 11.94, and CO: 0.63.

For the city of Chennai:

- The average AQI is 114.50, falling into the 'Satisfactory' category.
- Key pollutant levels on average are: PM2.5: 50.21, PM10: 109.26, NO₂: 16.59, and CO: 1.02.
- PM10 levels are also elevated, contributing to reduced air quality.

For the city of Coimbatore:

- The average AQI is 73.02, falling into the 'Satisfactory' category.
- Key pollutant levels on average are: PM2.5: 29.55, PM10: 39.13, NO₂: 27.68, and CO: 0.95.

For the city of Delhi:

- The average AQI is 259.49, falling into the 'Poor' category.
- Key pollutant levels on average are: PM2.5: 117.52, PM10: 228.94, NO₂: 50.75, and CO: 1.99.
- PM2.5 levels are relatively high, which may pose a health risk.
- PM10 levels are also elevated, contributing to reduced air quality.
- NO₂ levels exceed safe limits, indicating a concern for air pollution from combustion sources.

For the city of Ernakulam:

- The average AQI is 92.36, falling into the 'Satisfactory' category.
- Key pollutant levels on average are: PM2.5: 25.16, PM10: 48.40, NO₂: 12.12, and CO: 1.63.

For the city of Gurugram:

- The average AQI is 225.12, falling into the 'Moderate' category.
- Key pollutant levels on average are: PM2.5: 115.44, PM10: 155.23, NO₂: 24.04, and CO: 1.34.
- PM2.5 levels are relatively high, which may pose a health risk.
- PM10 levels are also elevated, contributing to reduced air quality.

For the city of Guwahati:

- The average AQI is 140.11, falling into the 'Satisfactory' category.
- Key pollutant levels on average are: PM2.5: 61.50, PM10: 113.82, NO₂: 13.54, and CO: 0.73.
- PM2.5 levels are relatively high, which may pose a health risk.
- PM10 levels are also elevated, contributing to reduced air quality.

For the city of Hyderabad:

- The average AQI is 109.21, falling into the 'Moderate' category.
- Key pollutant levels on average are: PM2.5: 47.15, PM10: 95.34, NO₂: 28.50, and CO: 0.61.

For the city of Jaipur:

- The average AQI is 133.68, falling into the 'Moderate' category.
- Key pollutant levels on average are: PM2.5: 54.71, PM10: 123.78,

NO_2 : 32.46, and CO : 0.81.

- PM_{10} levels are also elevated, contributing to reduced air quality.

For the city of Jorapokhar:

- The average AQI is 159.25, falling into the 'Moderate' category.

- Key pollutant levels on average are: $\text{PM}_{2.5}$: 66.02, PM_{10} : 151.88, NO_2 : 9.58, and CO : 1.20.

- $\text{PM}_{2.5}$ levels are relatively high, which may pose a health risk.

- PM_{10} levels are also elevated, contributing to reduced air quality.

For the city of Kochi:

- The average AQI is 104.28, falling into the 'Satisfactory' category.

- Key pollutant levels on average are: $\text{PM}_{2.5}$: 31.54, PM_{10} : 66.11, NO_2 : 15.00, and CO : 1.29.

For the city of Kolkata:

- The average AQI is 140.57, falling into the 'Satisfactory' category.

- Key pollutant levels on average are: $\text{PM}_{2.5}$: 64.65, PM_{10} : 116.08, NO_2 : 40.80, and CO : 0.85.

- $\text{PM}_{2.5}$ levels are relatively high, which may pose a health risk.

- PM_{10} levels are also elevated, contributing to reduced air quality.

- NO_2 levels exceed safe limits, indicating a concern for air pollution from combustion sources.

For the city of Lucknow:

- The average AQI is 217.97, falling into the 'Moderate' category.

- Key pollutant levels on average are: $\text{PM}_{2.5}$: 109.75, PM_{10} : 118.13, NO_2 : 33.88, and CO : 1.75.

- $\text{PM}_{2.5}$ levels are relatively high, which may pose a health risk.

- PM_{10} levels are also elevated, contributing to reduced air quality.

For the city of Mumbai:

- The average AQI is 105.35, falling into the 'Satisfactory' category.

- Key pollutant levels on average are: $\text{PM}_{2.5}$: 35.30, PM_{10} : 97.24, NO_2 : 25.59, and CO : 1.23.

For the city of Patna:

- The average AQI is 240.78, falling into the 'Moderate' category.

- Key pollutant levels on average are: $\text{PM}_{2.5}$: 124.65, PM_{10} : 119.26, NO_2 : 39.06, and CO : 1.59.

- $\text{PM}_{2.5}$ levels are relatively high, which may pose a health risk.

- PM_{10} levels are also elevated, contributing to reduced air quality.

For the city of Shillong:

- The average AQI is 53.80, falling into the 'Good' category.

- Key pollutant levels on average are: $\text{PM}_{2.5}$: 25.77, PM_{10} : 35.84, NO_2 : 3.50, and CO : 0.27.

For the city of Talcher:

- The average AQI is 172.89, falling into the 'Moderate' category.


- Key pollutant levels on average are: PM2.5: 62.25, PM10: 165.72, NO2: 15.24, and CO: 1.82.
- PM2.5 levels are relatively high, which may pose a health risk.
- PM10 levels are also elevated, contributing to reduced air quality.

For the city of Thiruvananthapuram:

- The average AQI is 75.88, falling into the 'Satisfactory' category.
- Key pollutant levels on average are: PM2.5: 28.61, PM10: 53.59, NO2: 9.54, and CO: 0.98.


For the city of Visakhapatnam:

- The average AQI is 117.27, falling into the 'Moderate' category.
- Key pollutant levels on average are: PM2.5: 47.75, PM10: 107.02, NO2: 37.10, and CO: 0.86.
- PM10 levels are also elevated, contributing to reduced air quality.



```
C:\Users\kadam\AppData\Local\Temp\ipykernel_11664\633187323.py:62:
UserWarning: Creating legend with loc="best" can be slow with large
amounts of data.
```

```
plt.tight_layout()
```




```
plt.hist(cleandt[], histtype='bar', rwidth=0.8)
plt.xlabel('age groups')
plt.ylabel('Number of people')
plt.title('Histogram')
plt.show()

Cell In[14], line 1
    plt.hist(cleandt[], histtype='bar', rwidth=0.8)
    ^
SyntaxError: invalid syntax. Perhaps you forgot a comma?
```

```
aqi_count= cleandt.AQI_Bucket.value_counts()
aqi_count
```


```
AQI_Bucket
Moderate      8829
Satisfactory  8224
Poor          2781
Very Poor     2337
Good          1341
Severe         1338
Name: count, dtype: int64
```

```
# Violin plot for AQI by AQI_Bucket
sns.violinplot(x='AQI_Bucket', y='AQI', data=cleandt)
plt.title('Distribution of AQI by AQI Bucket')
plt.show()
```




```
# Correlation matrix for numerical variables
correlation_matrix = cleandt[['PM2.5', 'PM10', 'N02', 'S02',
'AQI']].corr()


# Plotting heatmap
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap of Air Quality Indicators')
plt.show()
```



```
sns.boxplot(x='AQI_Bucket', y='AQI', data=cleandt)
plt.title('Distribution of AQI by AQI Bucket')
plt.show()
```



```
plt.figure(figsize=(12,6))
# plt.title(schema.Gender)
plt.pie(aqi_count, labels=aqi_count.index, autopct='%.1f%%',
startangle=180);
```



Linear Regression Model

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# creating 2 datasets.
# x contains all the columns essential to predict the AQI.
# y contains the AQI itself.
# we are going to predict the AQI based on air components.
x = cleandt.drop(['City','Date','AQI','AQI_Bucket'],axis=1)
y = cleandt['AQI']
x

```

	PM2.5	PM10	NO	NO2	NOx	NH3	CO
S02 \	83.13	118.127103	6.93	28.71	33.72	23.483476	6.93
49.52	79.84	118.127103	13.85	28.68	41.08	23.483476	13.85
48.49							

2	94.52	118.127103	24.39	32.66	52.61	23.483476	24.39
67.39							
3	135.99	118.127103	43.48	42.08	84.57	23.483476	43.48
75.23							
4	178.33	118.127103	54.56	35.31	72.80	23.483476	54.56
55.04							
...
.							
24845	15.02	50.940000	7.68	25.06	19.54	12.470000	0.47
8.55							
24846	24.38	74.090000	3.42	26.06	16.53	11.990000	0.52
12.72							
24847	22.91	65.730000	3.45	29.53	18.33	10.710000	0.48
8.42							
24848	16.64	49.970000	4.05	29.26	18.80	10.030000	0.52
9.84							
24849	15.00	66.000000	0.40	26.85	14.05	5.200000	0.59
2.10							

	O3	Benzene	Toluene	Xylene
0	59.76	0.02000	0.000000	3.140000
1	97.07	0.04000	0.000000	4.810000
2	111.33	0.24000	0.010000	7.670000
3	102.70	0.40000	0.040000	25.870000
4	107.38	0.46000	0.060000	35.610000
...
24845	23.30	2.24000	12.070000	0.730000
24846	30.14	0.74000	2.210000	0.380000
24847	30.96	0.01000	0.010000	0.000000
24848	28.30	0.00000	0.000000	0.000000
24849	17.05	3.28084	8.700972	3.070128

[24850 rows x 12 columns]

#splitting the dataset

```
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=42)
```

#model initialization

```
model = LinearRegression()
model.fit(x_train, y_train)
```

LinearRegression()

```
print(model.intercept_)
print(model.coef_)
```

6.768096690239048

```
[ 1.11978235  0.29197875 -0.1353422   0.24151263  0.09964323 -
0.0675648
```

```
11.72232249  0.68818773  0.1982416   0.02099309 -0.02508842 -
0.13344967]

#predicting the AQI
y_pred = model.predict(x_test)

print(y_pred)
[114.04836463 200.79703022 116.08024452 ... 172.82832023 109.77860936
120.67003006]


#rmse : Root Mean Squared Error
#RMSE is the square root of the average squared differences between
#the predicted values (y_pred) and the actual values (y_test)

#R2 (R-squared), or the coefficient of determination, is a statistical
#measure that tells how well the independent variables (features)
#explain the variance in the dependent variable (target).
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)

print(f"Root Mean Squared Error (RMSE): {rmse}")
print(f"R2 Score: {r2}")

Root Mean Squared Error (RMSE): 40.783195026017644
R2 Score: 0.9091656886933218

plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, alpha=0.5)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--')
plt.title("Actual vs Predicted AQI")
plt.xlabel("Actual AQI")
plt.ylabel("Predicted AQI")
plt.show()
```



```

import pandas as pd

cgr_data = {'category': ['Severe', 'Very Poor', 'Poor', 'Moderate',
'Satisfactory', 'Good']}
cgr_num = pd.DataFrame()
# Mapping dictionary
mapping = {'Severe': 0, 'Very Poor': 1, 'Poor': 2, 'Moderate': 3, 'Satisfactory': 4, 'Good': 5,}

# Creating a new column with mapped values
cleandt['AQI_Bucket_num'] = cleandt['AQI_Bucket'].map(mapping)
cleandt=cleandt.drop('numeric_category',axis=1)

cleandt

```

	City	Date	PM2.5	PM10	NO	NO2
N0x \						
0	Ahmedabad	2015-01-29	83.13	118.127103	6.93	28.71
33.72						
1	Ahmedabad	2015-01-30	79.84	118.127103	13.85	28.68
41.08						

2	Ahmedabad	2015-01-31	94.52	118.127103	24.39	32.66	
52.61							
3	Ahmedabad	2015-02-01	135.99	118.127103	43.48	42.08	
84.57							
4	Ahmedabad	2015-02-02	178.33	118.127103	54.56	35.31	
72.80							
...	
...							
24845	Visakhapatnam	2020-06-27	15.02	50.940000	7.68	25.06	
19.54							
24846	Visakhapatnam	2020-06-28	24.38	74.090000	3.42	26.06	
16.53							
24847	Visakhapatnam	2020-06-29	22.91	65.730000	3.45	29.53	
18.33							
24848	Visakhapatnam	2020-06-30	16.64	49.970000	4.05	29.26	
18.80							
24849	Visakhapatnam	2020-07-01	15.00	66.000000	0.40	26.85	
14.05							
AQI \	NH3	CO	S02	O3	Benzene	Toluene	Xylene
0	23.483476	6.93	49.52	59.76	0.02000	0.000000	3.140000
209.0							
1	23.483476	13.85	48.49	97.07	0.04000	0.000000	4.810000
328.0							
2	23.483476	24.39	67.39	111.33	0.24000	0.010000	7.670000
514.0							
3	23.483476	43.48	75.23	102.70	0.40000	0.040000	25.870000
782.0							
4	23.483476	54.56	55.04	107.38	0.46000	0.060000	35.610000
914.0							
...
...							
24845	12.470000	0.47	8.55	23.30	2.24000	12.070000	0.730000
41.0							
24846	11.990000	0.52	12.72	30.14	0.74000	2.210000	0.380000
70.0							
24847	10.710000	0.48	8.42	30.96	0.01000	0.010000	0.000000
68.0							
24848	10.030000	0.52	9.84	28.30	0.00000	0.000000	0.000000
54.0							
24849	5.200000	0.59	2.10	17.05	3.28084	8.700972	3.070128
50.0							
	AQI_Bucket	AQI_Bucket_num					
0	Poor	2					
1	Very Poor	1					
2	Severe	0					
3	Severe	0					

```

4           Severe          0
...
24845      Good           5
24846  Satisfactory     4
24847  Satisfactory     4
24848  Satisfactory     4
24849      Good           5

[24850 rows x 17 columns]

a = cleandt.drop(['City', 'Date', 'AQI_Bucket_num', 'AQI_Bucket'], axis=1)
b = cleandt['AQI_Bucket_num']
a

    PM2.5      PM10       NO      NO2      NOx      NH3       CO
S02 \
0    83.13  118.127103   6.93   28.71   33.72  23.483476   6.93
49.52
1    79.84  118.127103  13.85   28.68   41.08  23.483476  13.85
48.49
2    94.52  118.127103  24.39   32.66   52.61  23.483476  24.39
67.39
3   135.99  118.127103  43.48   42.08   84.57  23.483476  43.48
75.23
4   178.33  118.127103  54.56   35.31   72.80  23.483476  54.56
55.04
...
.
24845  15.02  50.940000   7.68   25.06   19.54  12.470000   0.47
8.55
24846  24.38  74.090000   3.42   26.06   16.53  11.990000   0.52
12.72
24847  22.91  65.730000   3.45   29.53   18.33  10.710000   0.48
8.42
24848  16.64  49.970000   4.05   29.26   18.80  10.030000   0.52
9.84
24849  15.00  66.000000   0.40   26.85   14.05   5.200000   0.59
2.10

    O3  Benzene      Toluene      Xylene      AQI
0    59.76  0.02000  0.000000  3.140000  209.0
1    97.07  0.04000  0.000000  4.810000  328.0
2   111.33  0.24000  0.010000  7.670000  514.0
3   102.70  0.40000  0.040000  25.870000  782.0
4   107.38  0.46000  0.060000  35.610000  914.0
...
24845  23.30  2.24000  12.070000  0.730000   41.0
24846  30.14  0.74000  2.210000  0.380000   70.0
24847  30.96  0.01000  0.010000  0.000000   68.0
24848  28.30  0.00000  0.000000  0.000000   54.0

```

```
24849    17.05   3.28084    8.700972    3.070128    50.0
[24850 rows x 13 columns]
b
0      2
1      1
2      0
3      0
4      0
..
24845    5
24846    4
24847    4
24848    4
24849    5
Name: AQI_Bucket_num, Length: 24850, dtype: int64

a_train, a_test, b_train, b_test = train_test_split(a, b, test_size =
0.20, random_state = 42, stratify = cleandt['AQI_Bucket_num'])

from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(a_train, b_train)
b_predict = model.predict(a_test)

C:\Users\kadam\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as
shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear\_model.html#logistic-
regression
n_iter_i = _check_optimize_result(
b_predict
array([4, 1, 3, ..., 3, 4, 1])


# Confusion matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(b_test, b_predict)
cm

array([[ 163,    78,    12,    15,      0,      0],
       [  57,   255,    71,    82,      2,      0],
```

```
[ 22, 130, 169, 229, 6, 0],  
[ 19, 74, 124, 1306, 238, 5],  
[ 15, 22, 12, 849, 728, 19],  
[ 0, 1, 0, 37, 196, 34]])
```

```
sns.heatmap(pd.DataFrame(cm), annot=True)
```

```
<Axes: >
```



```
from sklearn.metrics import accuracy_score  
  
accuracy =accuracy_score(b_test, b_predict)  
accuracy  
  
0.53420523138833  
  
b_predict =  
model.predict([[1,148,72,35,79.799,33.6,0.627,50,0,3,5,4,2]])  
print(b_predict)  
if b_predict==0:  
    print("Severe")  
else:  
    print("Not Severe")  
  
[5]  
Not Severe
```

```

C:\Users\kadam\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\base.py:493: UserWarning: X does not have valid
feature names, but LogisticRegression was fitted with feature names
warnings.warn(
import pandas as pd

# Step 2: Create a dictionary that maps each city to its state
city_to_state = {
    "Ahmedabad": "Gujarat",
    "Bengaluru": "Karnataka",
    "Chennai": "Tamil Nadu",
    "Mumbai": "Maharashtra",
    "Lucknow": "Uttar Pradesh",
    "Delhi": "Delhi",
    "Hyderabad": "Telangana",
    "Patna": "Bihar",
    "Gurugram": "Haryana",
    "Visakhapatnam": "Andhra Pradesh",
    "Amritsar": "Punjab",
    "Jorapokhar": "Jharkhand",
    "Jaipur": "Rajasthan",
    "Thiruvananthapuram": "Kerala",
    "Amaravati": "Andhra Pradesh",
    "Brajrajnagar": "Odisha",
    "Talcher": "Odisha",
    "Kolkata": "West Bengal",
    "Guwahati": "Assam",
    "Coimbatore": "Tamil Nadu",
    "Shillong": "Meghalaya",
    "Chandigarh": "Chandigarh",
    "Bhopal": "Madhya Pradesh",
    "Kochi": "Kerala",
    "Ernakulam": "Kerala",
    "Aizawl": "Mizoram"
}
# Step 3: Add a new column 'State' by mapping the cities to their
states
cleandt['State'] = cleandt['City'].map(city_to_state)

# Display the DataFrame
print(cleandt)

      City        Date   PM2.5      PM10      NO     NO2
NOx \
0      Ahmedabad  2015-01-29  83.13  118.127103  6.93  28.71
33.72
1      Ahmedabad  2015-01-30  79.84  118.127103 13.85  28.68
41.08

```

2	Ahmedabad	2015-01-31	94.52	118.127103	24.39	32.66	
52.61							
3	Ahmedabad	2015-02-01	135.99	118.127103	43.48	42.08	
84.57							
4	Ahmedabad	2015-02-02	178.33	118.127103	54.56	35.31	
72.80							
...	
...							
24845	Visakhapatnam	2020-06-27	15.02	50.940000	7.68	25.06	
19.54							
24846	Visakhapatnam	2020-06-28	24.38	74.090000	3.42	26.06	
16.53							
24847	Visakhapatnam	2020-06-29	22.91	65.730000	3.45	29.53	
18.33							
24848	Visakhapatnam	2020-06-30	16.64	49.970000	4.05	29.26	
18.80							
24849	Visakhapatnam	2020-07-01	15.00	66.000000	0.40	26.85	
14.05							
AQI \	NH3	CO	S02	O3	Benzene	Toluene	Xylene
0	23.483476	6.93	49.52	59.76	0.02000	0.000000	3.140000
209.0							
1	23.483476	13.85	48.49	97.07	0.04000	0.000000	4.810000
328.0							
2	23.483476	24.39	67.39	111.33	0.24000	0.010000	7.670000
514.0							
3	23.483476	43.48	75.23	102.70	0.40000	0.040000	25.870000
782.0							
4	23.483476	54.56	55.04	107.38	0.46000	0.060000	35.610000
914.0							
...
...							
24845	12.470000	0.47	8.55	23.30	2.24000	12.070000	0.730000
41.0							
24846	11.990000	0.52	12.72	30.14	0.74000	2.210000	0.380000
70.0							
24847	10.710000	0.48	8.42	30.96	0.01000	0.010000	0.000000
68.0							
24848	10.030000	0.52	9.84	28.30	0.00000	0.000000	0.000000
54.0							
24849	5.200000	0.59	2.10	17.05	3.28084	8.700972	3.070128
50.0							
	AQI_Bucket			State			
0	Poor			Gujarat			
1	Very Poor			Gujarat			
2	Severe			Gujarat			
3	Severe			Gujarat			

4	Severe	Gujarat
...
24845	Good	Andhra Pradesh
24846	Satisfactory	Andhra Pradesh
24847	Satisfactory	Andhra Pradesh
24848	Satisfactory	Andhra Pradesh
24849	Good	Andhra Pradesh

[24850 rows x 17 columns]

cleandt

NOx	City	Date	PM2.5	PM10	NO	NO2
0	Ahmedabad	2015-01-29	83.13	118.127103	6.93	28.71
33.72	Ahmedabad	2015-01-30	79.84	118.127103	13.85	28.68
41.08	Ahmedabad	2015-01-31	94.52	118.127103	24.39	32.66
52.61	Ahmedabad	2015-02-01	135.99	118.127103	43.48	42.08
84.57	Ahmedabad	2015-02-02	178.33	118.127103	54.56	35.31
72.80
24845	Visakhapatnam	2020-06-27	15.02	50.940000	7.68	25.06
19.54	Visakhapatnam	2020-06-28	24.38	74.090000	3.42	26.06
16.53	Visakhapatnam	2020-06-29	22.91	65.730000	3.45	29.53
18.33	Visakhapatnam	2020-06-30	16.64	49.970000	4.05	29.26
18.80	Visakhapatnam	2020-07-01	15.00	66.000000	0.40	26.85
14.05

```
...
24845 12.470000 0.47 8.55 23.30 2.24000 12.070000 0.730000
41.0
24846 11.990000 0.52 12.72 30.14 0.74000 2.210000 0.380000
70.0
24847 10.710000 0.48 8.42 30.96 0.01000 0.010000 0.000000
68.0
24848 10.030000 0.52 9.84 28.30 0.00000 0.000000 0.000000
54.0
24849 5.200000 0.59 2.10 17.05 3.28084 8.700972 3.070128
50.0
```

	AQI_Bucket	State
0	Poor	Gujarat
1	Very Poor	Gujarat
2	Severe	Gujarat
3	Severe	Gujarat
4	Severe	Gujarat
...
24845	Good	Andhra Pradesh
24846	Satisfactory	Andhra Pradesh
24847	Satisfactory	Andhra Pradesh
24848	Satisfactory	Andhra Pradesh
24849	Good	Andhra Pradesh

[24850 rows x 17 columns]

```
state_aqi_avg = cleandt.groupby("State")["AQI"].mean()

state_aqi_avg
```

State	
Andhra Pradesh	108.086481
Assam	140.111111
Bihar	240.782042
Chandigarh	96.498328
Delhi	259.487744
Gujarat	452.122939
Haryana	225.123882
Jharkhand	159.251621
Karnataka	94.318325
Kerala	81.021277
Madhya Pradesh	132.827338
Maharashtra	105.352258
Meghalaya	53.795122
Mizoram	34.765766
Odisha	161.463501
Punjab	119.920959
Rajasthan	133.679159
Tamil Nadu	108.098294

```

Telangana      109.207447
Uttar Pradesh  217.973059
West Bengal    140.566313
Name: AQI, dtype: float64

import pandas as pd


# Convert the 'Date' column to datetime format
cleandt['Date'] = pd.to_datetime(cleandt['Date'])

# Extract the year from the 'Date' column
cleandt['Year'] = cleandt['Date'].dt.year

# Group by 'Year' and calculate the average AQI for each year
yearly_aqi_avg15 = cleandt[cleandt['Date'].dt.year ==
                           2015].groupby("State")["AQI"].mean().reset_index()
yearly_aqi_avg = cleandt.groupby(["Year", "State"])["AQI"].mean().reset_index()

# Display the result
#print("Average AQI by Year:")
yearly_aqi_avg

   Year        State     AQI
0  2015       Bihar  350.555556
1  2015       Delhi  297.024658
2  2015     Gujarat  310.950570
3  2015   Karnataka  112.573427
4  2015  Tamil Nadu  148.333333
..  ...
82 2020  Rajasthan  105.120219
83 2020  Tamil Nadu  74.705015
84 2020  Telangana  78.174863
85 2020  Uttar Pradesh  157.125683
86 2020  West Bengal  117.295082

[87 rows x 3 columns]

yearly_aqi_avg16 = cleandt[cleandt['Date'].dt.year ==
                           2015].groupby("State")["AQI"].mean().reset_index()
yearly_aqi_avg16

   State     AQI
0  Bihar  350.555556
1  Delhi  297.024658
2  Gujarat  310.950570
3  Karnataka  112.573427
4  Tamil Nadu  148.333333

```

```

5      Telangana  143.419118
6  Uttar Pradesh  202.235915

yearly_aqi_avg17 = cleandt[cleandt['Date'].dt.year ==
2015].groupby("State")["AQI"].mean().reset_index()
yearly_aqi_avg17

      State      AQI
0    Bihar  350.555556
1    Delhi  297.024658
2  Gujarat  310.950570
3  Karnataka  112.573427
4  Tamil Nadu  148.333333
5   Telangana  143.419118
6  Uttar Pradesh  202.235915

yearly_aqi_avg18 = cleandt[cleandt['Date'].dt.year ==
2015].groupby("State")["AQI"].mean().reset_index()
yearly_aqi_avg18

      State      AQI
0    Bihar  350.555556
1    Delhi  297.024658
2  Gujarat  310.950570
3  Karnataka  112.573427
4  Tamil Nadu  148.333333
5   Telangana  143.419118
6  Uttar Pradesh  202.235915

yearly_aqi_avg19 = cleandt[cleandt['Date'].dt.year ==
2015].groupby("State")["AQI"].mean().reset_index()
yearly_aqi_avg19

      State      AQI
0    Bihar  350.555556
1    Delhi  297.024658
2  Gujarat  310.950570
3  Karnataka  112.573427
4  Tamil Nadu  148.333333
5   Telangana  143.419118
6  Uttar Pradesh  202.235915

yearly_aqi_avg20 = cleandt[cleandt['Date'].dt.year ==
2015].groupby("State")["AQI"].mean().reset_index()
yearly_aqi_avg20

      State      AQI
0    Bihar  350.555556
1    Delhi  297.024658
2  Gujarat  310.950570
3  Karnataka  112.573427

```

```

4      Tamil Nadu  148.333333
5      Telangana   143.419118
6  Uttar Pradesh  202.235915

# Convert Date to datetime
cleandt['Date'] = pd.to_datetime(cleandt['Date'])

# Initial data exploration
print(cleandt.info())
print(cleandt.describe())


# Drop any missing values or fill them as appropriate
cleandt = cleandt.dropna() # Adjust this step as necessary for your
                           data

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24850 entries, 0 to 24849
Data columns (total 16 columns):
 #   Column      Non-Null Count Dtype  
 --- 
 0   City        24850 non-null  object  
 1   Date        24850 non-null  datetime64[ns]
 2   PM2.5       24850 non-null  float64 
 3   PM10        24850 non-null  float64 
 4   NO          24850 non-null  float64 
 5   NO2         24850 non-null  float64 
 6   NOx         24850 non-null  float64 
 7   NH3         24850 non-null  float64 
 8   CO          24850 non-null  float64 
 9   SO2         24850 non-null  float64 
 10  O3          24850 non-null  float64 
 11  Benzene     24850 non-null  float64 
 12  Toluene     24850 non-null  float64 
 13  Xylene      24850 non-null  float64 
 14  AQI         24850 non-null  float64 
 15  AQI_Bucket  24850 non-null  object  
dtypes: datetime64[ns](1), float64(13), object(2)
memory usage: 3.0+ MB
None
      Date      PM2.5      PM10 \
count  24850  24850.000000  24850.000000
mean   2018-07-24 18:51:25.714285568  67.475903  118.361096
min    2015-01-01 00:00:00            0.040000  0.030000
25%   2017-08-16 00:00:00            29.560000  71.780000
50%   2018-11-05 00:00:00            50.165000  118.127103
75%   2019-10-11 00:00:00            79.507500  122.957500
max   2020-07-01 00:00:00            914.940000  917.080000
std    NaN                62.208948  75.660501

```


	NO	NO2	NOx	NH3
C0 \				
count	24850.000000	24850.000000	24850.000000	24850.000000
24850.000000				
mean	17.621678	28.971818	32.290515	23.752394
2.343536				
min	0.030000	0.010000	0.000000	0.010000
0.000000				
25%	5.720000	12.090000	14.030000	11.280000
0.590000				
50%	10.075000	22.535000	25.720000	23.483476
0.950000				
75%	19.710000	37.910000	38.170000	24.710000
1.530000				
max	390.680000	362.210000	378.240000	352.890000
175.810000				
std	22.245860	24.432587	29.542968	22.214343
7.011582				
	S02	O3	Benzene	Toluene
Xylene \				
count	24850.000000	24850.000000	24850.000000	24850.000000
24850.000000				
mean	14.367049	34.899199	3.433371	9.332356
3.267909				
min	0.010000	0.010000	0.000000	0.000000
0.000000				
25%	5.790000	19.640000	0.340000	1.580000
2.650000				
50%	9.430000	32.060000	1.810000	6.790000
3.070128				
75%	14.890000	45.397500	3.280840	8.700972
3.070128				
max	186.080000	257.730000	455.030000	454.850000
170.370000				
std	17.215237	21.368979	14.851776	18.273322
4.178816				
	AQI			
count	24850.000000			
mean	166.463581			
min	13.000000			
25%	81.000000			
50%	118.000000			
75%	208.000000			
max	2049.000000			
std	140.696585			
# Set Date as index for time series analysis				
#cleandt.set_index('Date', inplace=True)				

```
# Example: Plot AQI trends over time for each city
plt.figure(figsize=(12, 6))
sns.lineplot(data=cleandt, x=cleandt.index, y="AQI", hue="City")
plt.title("AQI Trend Over Time by City")
plt.xlabel("Date")
plt.ylabel("AQI Level")
plt.legend(title="City")
plt.show()
```



```
# Seasonal patterns for a particular city
#city_data = cleandt[cleandt['City'] == 'CityName'] # Replace with
#actual city name
plt.figure(figsize=(10, 6))
sns.lineplot(x=cleandt.index, y=cleandt['AQI'])

plt.title("Seasonal AQI Variation for CityName")
plt.xlabel("Date")
plt.ylabel("AQI Level")
plt.show()
```




```
# Calculate average AQI by city
avg_aqi_by_city = data.groupby('City')['AQI'].mean().sort_values()
```

```
# Plot average AQI by city
plt.figure(figsize=(10, 6))
sns.barplot(x=avg_aqi_by_city.index, y=avg_aqi_by_city.values,
palette="viridis")
plt.xticks(rotation=90)
plt.title("Average AQI by City")
plt.xlabel("City")
plt.ylabel("Average AQI")
plt.show()
```

```
C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\600190197.py:6:
FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```


```
sns.barplot(x=avg_aqi_by_city.index, y=avg_aqi_by_city.values,
palette="viridis")
```



```
# Calculate distribution of AQI Buckets across cities
aqi_bucket_distribution = data.groupby(['City',
'AQI_Bucket']).size().unstack().fillna(0)

# Plot AQI Bucket distribution by city
aqi_bucket_distribution.plot(kind='bar', stacked=True, figsize=(14,
8), colormap="Set3")
plt.title("Distribution of AQI Buckets by City")
plt.xlabel("City")
plt.ylabel("Frequency")
plt.legend(title="AQI Bucket", bbox_to_anchor=(1.05, 1), loc='upper
left')
plt.show()

# Health impact assessment - percentage of "Unhealthy" AQI levels
unhealthy_percentage = (data['AQI_Bucket'] == 'Unhealthy').mean() *
100
print(f"Percentage of Unhealthy AQI Levels: {unhealthy_percentage:.2f}%
")
```



Percentage of Unhealthy AQI Levels: 0.00%

```
# Select features and target for prediction
# = cleandt[] # Independent variables
from sklearn.ensemble import RandomForestRegressor
X = cleandt.drop(['City','AQI','AQI_Bucket'],axis=1)
y = cleandt['AQI']           # Target variable

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Model: Random Forest Regressor
model = RandomForestRegressor(random_state=42)
model.fit(X_train, y_train)

# Prediction
y_pred = model.predict(X_test)

# Model evaluation
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")


# Plot actual vs predicted AQI
plt.figure(figsize=(8, 6))
```

```

plt.scatter(y_test, y_pred, alpha=0.5)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--')
plt.title("Actual vs Predicted AQI")
plt.xlabel("Actual AQI")
plt.ylabel("Predicted AQI")
plt.show()

```

Mean Squared Error: 1663.2689965301904
R-squared: 0.9091656886933218



```

# Calculate average AQI by city
avg_aqi_by_city = cleandt.groupby('City')['AQI'].mean().sort_values()

# Plot average AQI by city
plt.figure(figsize=(10, 6))
sns.barplot(x=avg_aqi_by_city.index, y=avg_aqi_by_city.values,
palette="viridis")
plt.xticks(rotation=90)
plt.title("Average AQI by City")
plt.xlabel("City")

```

```


plt.ylabel("Average AQI")
plt.show()

C:\Users\kadam\AppData\Local\Temp\ipykernel_21312\998449420.py:6:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.

sns.barplot(x=avg_aqi_by_city.index, y=avg_aqi_by_city.values,
palette="viridis")

```



```

# Correlation heatmap between pollutants and AQI
pollutants = ['PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2',
'03', 'Benzene', 'Toluene', 'Xylene']
plt.figure(figsize=(12, 8))
sns.heatmap(cleandt[pollutants + ['AQI']].corr(), annot=True,
cmap="coolwarm", vmin=-1, vmax=1)


```

```

plt.title("Correlation Between Pollutants and AQI")
plt.show()

# Correlation between AQI and AQI_Bucket (health impact)
data['AQI_Bucket'] = data['AQI_Bucket'].astype('category').cat.codes
# Convert AQI Bucket to numeric codes
plt.figure(figsize=(8, 6))
sns.scatterplot(data=cleandt, x="AQI", y="AQI_Bucket")
plt.title("Correlation Between AQI and AQI Bucket")
plt.xlabel("AQI")
plt.ylabel("AQI Bucket")
plt.show()

```



Correlation Between AQI and AQI Bucket

